

Reproducing “Experimental Analysis of Dense vs. Sparse Retrieval”

Gabriele Righi

University of Pisa

February 11, 2026



Introduction & Research Goals

Context: The “Operational” Gap

- Modern IR has shifted from Sparse (BM25) to Dense (Embeddings).
- Dense retrieval is effective but resource-intensive.

The 3 Research Questions (Lin et al., 2024)

This project reproduces the analysis to answer three key operational questions:

- ① **RQ1:** Is the complexity of HNSW always necessary, or is **Flat index** sufficient for typical workloads?
- ② **RQ2:** Can we use **INT8 Quantization** to save RAM without hurting retrieval quality (nDCG)?
- ③ **RQ3:** How does **Sparse Retrieval** (BM25/SPLADE) compare to Dense methods in a modern GPU environment?

Experimental Setup

- **Benchmark:** BEIR (Heterogeneous Evaluation).
- **Datasets Analyzed:** scifact, nfcorpus, cqadupstack (Android, Gaming, GIS).
- **Hardware:** NVIDIA T4 GPU (16GB VRAM).
 - *Constraints simulated a realistic production environment.*

Metrics

nDCG@10 Retrieval Quality (Ranking accuracy).

Recall@10 Retrieval Coverage (Relevant items found).

QPS Queries Per Second (Throughput/Speed).

Challenge I: The Missing Link in Pyserini

Original Plan: Use Pyserini (Lucene wrapper) for *both* Sparse (BM25) and Dense (HNSW) indexing, matching the paper exactly.

The Obstacle:

- Pyserini supports HNSW via Java, but the **Python API for adding vectors to HNSW indexes is limited/missing**.
- Impossible to feed custom BGE embeddings directly into Pyserini's HNSW from Python easily.

The Engineering Pivot: Hybrid Pipeline

- **Sparse Path:** Retained Pyserini for BM25 (Industry Standard).
- **Dense Path:** Migrated to **Faiss** (Facebook AI Similarity Search).
- **Why?** Faiss offers native **GPU support** and flexible python bindings for custom embeddings, enabling a fairer comparison for RQ1 (Flat vs HNSW).

Challenge II: The SPLADE Underperformance Mystery

The Anomaly While BM25 and Dense Retrieval matched the paper's baselines immediately, the standard SPLADE implementation (via Pyserini) failed.

| Method | Implementation | nDCG@10 | Status |
|---------------|------------------------|--------------|--------------------|
| BM25 | Pyserini (Lucene) | 0.323 | Success |
| BGE Dense | Faiss | 0.370 | Success |
| SPLADE | Pyserini Impact | 0.230 | Fail (-28%) |

Root Cause Analysis

- *Hypothesis A (Wrong Model)*: Switched to `selfdistil`. No change.
- *Hypothesis B (Quantization)*: Forced integer scaling ($w \times 100$). No change.
- **Conclusion:** Pyserini's "Black Box" Impact Indexing was introducing **lossy compression artifacts**, destroying the precision of sparse weights.

The Solution: Custom Matrix Engine

The Fix: Bypassed Pyserini. Built a **pure Python sparse engine** using SciPy.

Methodology

- **Direct Encoding:** Used HuggingFace to avoid quantization artifacts.
- **Vectorized Search:** Replaced index lookup with Matrix Multiplication.


Core Logic

$$S = D \times q^T$$

Impact

- **Quality:** nDCG@10 restored to 0.357 (vs 0.230).
- **Speed:** ~296 QPS (Efficient on Python).

Implementation Pipeline: The “Dual-Engine” Architecture



Architectural Highlights:

- **Parallel Processing:** Simultaneous execution of Dense and Sparse pipelines.
- **Custom Matrix Engine:** Replaced Pyserini’s “Black Box” indexer with a transparent vector-matrix multiplication engine for SPLADE.
- **GPU Acceleration:** Integrated Faiss to unlock GPU speeds for Dense Retrieval.

Operational Advice: RQ1 & RQ2

Summary: When to switch Index? When to Quantize?

| Dataset | Method | Build | RAM | QPS | nDCG | Verdict / Trade-off |
|------------------------|-------------|-------|--------|-------|-------|---------------------|
| SciFact (5k) | Flat (FP32) | <0.1s | .01 GB | 4,205 | 0.738 | Best Choice |
| | Flat (INT8) | <0.1s | .01 GB | 2,209 | 0.736 | Slower (CPU limit) |
| Quora (523k) | Flat (FP32) | 1s | .38 GB | 120 | 0.889 | Too Slow |
| | HNSW (FP32) | 4m | .45 GB | 361 | 0.889 | Worth the wait |
| NQ (2.6M) | HNSW (FP32) | 19m | 2.4 GB | 9,034 | 0.464 | High RAM |
| | HNSW (INT8) | 25m | 0.6 GB | 9,100 | 0.468 | -75% RAM (Safe) |

Final Recommendations

- RQ1 (Scale):** Use **Flat** for small data (<100k). Use **HNSW** for large data (>200k) as QPS scales better, with negligible nDCG drop (-0.004).
- RQ2 (Quantization):** Always use **INT8** for HNSW. It saves 75% RAM while maintaining identical nDCG quality (0.541 vs 0.540).

Quality Analysis: Sparse vs. Learned Sparse vs. Dense

Q: Does SPLADE bridge the gap? Is Dense still superior?

| Dataset | BM25 | SPLADE | Dense | Winner |
|----------------------------------|-------|---------------|------------------------------------|--------------|
| SciFact <i>(Small)</i> | 0.679 | 0.717 | 0.738 <i>(Flat FP32)</i> | Dense |
| Quora <i>(Medium)</i> | 0.789 | 0.842 | 0.889 <i>(Flat FP32)</i> | Dense |
| NQ <i>(Large)</i> | 0.235 | 0.577* | 0.541 <i>(Flat FP32)</i> | Dense |

- **Small/Medium Data:** SPLADE improves over BM25 significantly (+15-20%).
- **NQ Result:** Dense is the winner, but required ~4h GPU encoding (vs BM25 minutes).
- ***Note on SPLADE:** Due to computational constraints, SPLADE was evaluated on a **100k subset**.
- **Conclusion:** Dense Retrieval provides the best quality, justifying the high setup cost.

Conclusions & Operational Advice

Answering the Research Questions:

RQ1 HNSW vs Flat: *Verdict:* HNSW is **not** always necessary. **Flat Index** is faster for datasets $< 100k$ docs and simpler to maintain.

RQ2 Quantization Safety: *Verdict:* Yes. INT8 reduces RAM by 75% with negligible quality loss ($< 1\%$ nDCG drop). It should be the default.

RQ3 Sparse vs Dense: *Verdict:* Dense methods (BGE) generally outperform Sparse (BM25) in quality, but custom Sparse implementations (Matrix Engine) can be extremely fast on GPU.

Final Recommendation

For most production scenarios under 1M docs: **Use Dense Retrieval with Flat Index + INT8 Quantization.**

References I

-  Formal, Thibault, Benjamin Piwowarski, and Stéphane Clinchant (2021). "SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking". In: *SIGIR '21*.
-  Lin, Sheng-Chieh, Shiguang Yang, and Jimmy Lin (2024). *Operational Advice for Dense and Sparse Retrievers: HNSW, Flat, or Inverted Indexes?* arXiv: 2409.06464 [cs.IR]. URL: <https://arxiv.org/abs/2409.06464>.
-  Righi, Gabriele (2026). *Dense vs Sparse Retrieval - Replication Repository*. GitHub Repository. URL: <https://github.com/Gheb6/dense-vs-sparse-retrieval>.
-  Robertson, Stephen and Hugo Zaragoza (2009). "The Probabilistic Relevance Framework: BM25 and Beyond". In: *Foundations and Trends® in Information Retrieval* 3.4, pp. 333–389. DOI: 10.1561/1500000019.

Thank you!

Questions?