

Reproducing “Experimental Analysis of Dense vs. Sparse Retrieval”

Gabriele Righi

University of Pisa

February 11, 2026



Introduction & Research Goals

Context: The “Operational” Gap

- Modern IR has shifted from Sparse (BM25) to Dense (Embeddings).
- Dense retrieval is effective but resource-intensive.

The 3 Research Questions (Lin et al., 2024)

This project reproduces the analysis to answer three key operational questions:

- ① **RQ1:** Is the complexity of HNSW always necessary, or is **Flat index** sufficient for typical workloads?
- ② **RQ2:** Can we use **INT8 Quantization** to save RAM without hurting retrieval quality (nDCG)?
- ③ **RQ3:** How does **Sparse Retrieval** (BM25/SPLADE) compare to Dense methods in a modern GPU environment?

Experimental Setup

- **Benchmark:** BEIR (Thakur et al. 2021) (Heterogeneous Evaluation).
- **Datasets Analyzed:** scifact, nfcorpus, cqadupstack (Android, Gaming, GIS).
- **Hardware:** NVIDIA T4 GPU (16GB VRAM).
 - *Constraints simulated a realistic production environment.*

Metrics

nDCG@10 Retrieval Quality (Ranking accuracy).

Recall@10 Retrieval Coverage (Relevant items found).

QPS Queries Per Second (Throughput/Speed).

Challenge I: The Missing Link in Pyserini

Original Plan: Use Pyserini (J. Lin et al. 2021) (Lucene wrapper) for *both* Sparse (BM25) and Dense (HNSW Malkov and Yashunin 2020) indexing, matching the paper exactly.

The Obstacle: Pyserini's **Python API for adding custom vectors to HNSW indexes is limited/missing**. Impossible to feed BGE embeddings into Pyserini's HNSW from Python.

The Engineering Pivot: Hybrid Pipeline

- **Sparse Path:** Retained Pyserini for BM25 (Industry Standard).
- **Dense Path:** Migrated to Faiss (Johnson, Douze, and Jégou 2021) (Facebook AI Similarity Search).

Why Faiss? Native **GPU support** and flexible Python bindings for custom embeddings, enabling a fair Flat vs HNSW comparison (RQ1).

Challenge II: The SPLADE Underperformance Mystery

The Anomaly While BM25 and Dense Retrieval matched the paper's baselines immediately, the standard SPLADE implementation (via Pyserini) failed.

Method	Implementation	nDCG@10	Status
BM25	Pyserini (Lucene)	0.323	Success
BGE Dense	Faiss	0.370	Success
SPLADE	Pyserini Impact	0.230	Fail (-28%)

Root Cause Analysis

- *Hypothesis A (Wrong Model)*: Switched to `selfdistil`. No change.
- *Hypothesis B (Quantization)*: Forced integer scaling ($w \times 100$). No change.
- **Conclusion:** Pyserini's "Black Box" Impact Indexing was introducing **lossy compression artifacts**, destroying the precision of sparse weights.

The Solution: Custom Matrix Engine

The Fix: Bypassed Pyserini. Built a **pure Python sparse engine** using SciPy.

Methodology

- **Direct Encoding:** Used HuggingFace to avoid quantization artifacts.
- **Vectorized Search:** Replaced index lookup with Matrix Multiplication.


Core Logic

$$S = D \times q^T$$

Impact

- **Quality:** nDCG@10 restored to 0.357 (vs 0.230).
- **Speed:** ~296 QPS (Efficient on Python).

Implementation Pipeline: The “Dual-Engine” Architecture



Architectural Highlights:

- **Parallel Processing:** Simultaneous execution of Dense and Sparse pipelines.
- **Custom Matrix Engine:** Replaced Pyserini’s “Black Box” indexer with a transparent vector-matrix multiplication engine for SPLADE.
- **GPU Acceleration:** Integrated Faiss to unlock GPU speeds for Dense Retrieval.

Operational Advice: RQ1 & RQ2

Summary: When to switch Index? When to Quantize?

Dataset	Method	Build	RAM	QPS	nDCG	Verdict / Trade-off
SciFact (5k)	Flat (FP32)	<0.1s	.01 GB	4,205	0.738	Best Choice
	Flat (INT8)	<0.1s	.01 GB	2,209	0.736	Slower (CPU limit)
Quora (523k)	Flat (FP32)	1s	.38 GB	120	0.889	Too Slow
	HNSW (FP32)	4m	.45 GB	361	0.889	Worth the wait
NQ (2.6M)	HNSW (FP32)	19m	2.4 GB	9,034	0.464	High RAM
	HNSW (INT8)	25m	0.6 GB	9,100	0.468	-75% RAM (Safe)

Final Recommendations

- RQ1 (Scale):** Use **Flat** for small data (<100k). Use **HNSW** for large data (>200k) as QPS scales better, with negligible nDCG drop (-0.004).
- RQ2 (Quantization):** Always use **INT8** for HNSW. It saves 75% RAM while maintaining identical nDCG quality (0.541 vs 0.540).

Quality Analysis: Sparse vs. Learned Sparse vs. Dense

Q: Does SPLADE bridge the gap? Is Dense still superior?

Dataset	BM25	SPLADE	Dense	Winner
SciFact <i>(Small)</i>	0.679	0.717	0.738 <i>(Flat FP32)</i>	Dense
Quora <i>(Medium)</i>	0.789	0.842	0.889 <i>(Flat FP32)</i>	Dense
NQ <i>(Large)</i>	0.235	0.577*	0.541 <i>(Flat FP32)</i>	Dense

- **Small/Medium Data:** SPLADE improves over BM25 significantly (+15-20%).
- **NQ Result:** Dense is the winner, but required ~4h GPU encoding (vs BM25 minutes).
- ***Note on SPLADE:** Due to computational constraints, SPLADE was evaluated on a **100k subset**.
- **Conclusion:** Dense Retrieval provides the best quality, justifying the high setup cost.

Conclusions & Operational Advice

Answering the Research Questions:

RQ1 HNSW vs Flat: *Verdict:* HNSW is **not** always necessary. **Flat Index** is faster for datasets $< 100k$ docs and simpler to maintain.

RQ2 Quantization Safety: *Verdict:* Yes. INT8 reduces RAM by 75% with negligible quality loss ($< 1\%$ nDCG drop). It should be the default.

RQ3 Sparse vs Dense: *Verdict:* Dense methods generally outperform Sparse (BM25) in quality, but custom Sparse implementations (Matrix Engine) can be extremely fast on GPU.

Final Recommendation

For most production scenarios under 1M docs: **Use Dense Retrieval with Flat Index + INT8 Quantization.**

References I

-  Formal, Thibault, Benjamin Piwowarski, and Stéphane Clinchant (2021). "SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking". In: *SIGIR '21*.
-  Johnson, Jeff, Matthijs Douze, and Hervé Jégou (2021). "Billion-scale similarity search with GPUs". In: *IEEE Transactions on Big Data* 7.3, pp. 535–547. DOI: 10.1109/TBDATA.2019.2921572.
-  Lin, Jimmy et al. (2021). "Pyserini: A Python Toolkit for Reproducible Information Retrieval Research with Sparse and Dense Representations". In: *SIGIR '21*, pp. 2356–2362. DOI: 10.1145/3404835.3463238.
-  Lin, Sheng-Chieh, Shiguang Yang, and Jimmy Lin (2024). *Operational Advice for Dense and Sparse Retrievers: HNSW, Flat, or Inverted Indexes?* arXiv: 2409.06464 [cs.IR]. URL: <https://arxiv.org/abs/2409.06464>.

References II

-  **Malkov, Yu A and D A Yashunin (2020).** "Efficient and robust approximate nearest neighbor search using Hierarchical Navigable Small World graphs". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 42.4, pp. 824–836. DOI: 10.1109/TPAMI.2018.2889473.
-  **Righi, Gabriele (2026).** *Dense vs Sparse Retrieval - Replication Repository*. GitHub Repository. URL: <https://github.com/Gheb6/dense-vs-sparse-retrieval>.
-  **Robertson, Stephen and Hugo Zaragoza (2009).** "The Probabilistic Relevance Framework: BM25 and Beyond". In: *Foundations and Trends® in Information Retrieval* 3.4, pp. 333–389. DOI: 10.1561/1500000019.
-  **Thakur, Nandan et al. (2021).** "BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models". In: *NeurIPS Datasets and Benchmarks*. URL: <https://arxiv.org/abs/2104.08663>.

Thank you!

Questions?